

Pattern inference theory: A probabilistic approach to vision*

Daniel Kersten and Paul Schrater[†]

August, 2001

*To be published: Kersten, D., & Schrater, P. R. Pattern inference theory: A probabilistic approach to vision. In Mausfeld, R., & Heyer, D. (Eds.), *Perception and the Physical World* (Chichester: John Wiley & Sons, Ltd.)

[†]Department of Psychology, University of Minnesota, Minneapolis, MN. Email: kersten@umn.edu.
Supported by NSF SBR-9631682 and NIH RO1 EY11507-001.

Abstract

The function of vision is to get correct and useful answers about the state of the world. However, given that the state of the world is not uniquely specified by the visual input, the visual system must make good guesses or *inferences*. Thus, theories of visual system functions will be theories of inference, and we need a language in which theories of inference can be described. Analogous to calculus having a minimum expressiveness required to formulate theories in physics, we argue that the language of Bayesian inference is fundamental to quantitatively describe how reliable answers about the world can be obtained from image patterns. Bayes provides a minimal formalism that can deal with the sophistication and versatility of perception missing from some other approaches. Key missing components include the ability to model uncertainty, probabilistic modeling of pattern synthesis as a necessary prerequisite to understanding pattern inference, the means to handle the complexity of natural images, and the diversity of visual tasks.

Most of the formal elements that we describe are not new and have their roots in signal detection theory and ideal observer analysis. We start from there to review and codify principles drawn from recent applications of Bayesian decision theory, Bayes nets and pattern theory to vision. To emphasize the importance of dealing with the complexity of natural image and scene patterns, we call the conjunction of principles drawn from these contributions *pattern inference theory*. Because of its generality, we do not see pattern inference theory as an experimentally testable theory of vision; however, it does provide a set of concepts and principles to formulate testable models. The test for a good theoretical framework is utility and completeness for deriving predictive theories. To illustrate the utility of the approach, we propose Bayesian principles of least commitment and modularity, each of which leads to testable hypotheses. Several recent examples of pattern inference theories are reviewed.

1 Perception is pattern decoding

Few would dispute the view that visual perception is the brain’s process for arriving at useful information about the world from images. Divergent opinions, however, have been expressed over how to describe the computations (or lack thereof) underlying visual behavior. Visual perception has been described as unconscious inference (Helmholtz and Southall, 1924; Gregory, 1980), reconstruction (Craik, 1943), resonance (Gibson, 1966), problem solving (Rock, 1983), computation (Marr, 1982), and more recently as Bayesian inference (Knill and Richards, 1996). In part, the debate gets muddled due to lack of a well-specified explanatory goal and level of abstraction. To clarify, we see the grand challenge to be the development of testable, quantitative theories of visual performance that take into account the complexities of natural images and the richness of visual behavior. But here the level of explanation is crucial: if our theories are too abstract, we lose the specificity of quantitative predictions; if the theories are too fine-grained, the model mechanisms for natural pattern processing will be too complex to test.

Our proposed strategy follows that of statistical mechanics. Few physicists doubt that the large-scale properties of physical systems rest on the lawful function of individual molecules, just as few brain scientists doubt that an organism’s behavior depends on the lawful function of neurons. Physicists would agree that the modeling level has to be appropriate to the measurements and phenomena of large-scale systems; thus statistical mechanics links molecular kinetics to thermodynamics. Although the bridge between neurons and system behavior has yet to be built, the language of Bayesian statistics

provides the level of description analogous to thermodynamics ¹. For vision, theories at this level are testable at the level of visual information and perceptual constraints ², and are less committal about representations, algorithms, or mechanisms.

The purpose of this chapter is to describe the fundamental principles of value in addressing the grand challenge. These principles constitute what we will refer to as *pattern inference theory*. The basic elements of pattern inference theory are not new and have their mathematical roots in communication and information theory (Shannon and Weaver, 1949), Bayesian decision theory (Berger, 1985), pattern theory (Grenander, 1996), and Bayes nets (Pearl, 1988). The refinement of the principles are derived from a history of applications to human vision in the domains of signal detection theory (Green and Swets, 1974), ideal observer analysis (Geisler, 1989; Schrater, 1998), Bayesian inference and decision theory (Kersten, 1990; Yuille and Blthoff, 1996), and pattern theory (Mumford, 1996; Yuille et al., 1998). “Pattern theory” was developed by Ulf Grenander to describe the mathematical study of complex natural patterns (Grenander, 1993, 1996; Mumford, 1996; Yuille et al., 1998). Central features of pattern theory are the importance of modeling pattern generation, and that natural pattern variation is characterized by four fundamental classes of deformations ³. Further, the generative model is seen as an essential part of inference (e.g. via flexible templates to fit incoming data in a feedback stage) to deal with certain types of deformation, such as occlusion (Mumford, 1994). Our particular emphasis is based on the synthesis and application of pattern theory and Bayesian decision theory to human vision (Yuille et al., 1998). As an elaboration of signal detection theory, we choose the words *pattern* and *inference* to stress the importance of modeling complex natural signals, and of considering tasks in addition to detection, respectively. We argue that pattern inference theory provides the best language for formulating quantitative theories of visual perception and action at the level of the naturally behaving (human) visual system.

Our goal is to derive probabilistic models of the observer’s world and sensory input, restricted by task. Such models have two components: the objects of the theory, and the operations of the theory. The objects of the theory are the set of possible image measurements I , the set of possible scene descriptions S , and the joint probability distribution of S and I : $p(S, I)$. The operations are given by the probability calculus, with decisions modeled as minimizing expected cost (or risk) given the probabilities. The richness of the theory lies in exploiting the structure induced in $p(S, I)$ by the regularities of the world (laws of physics) and by the habits of observers. A fundamental assumption of pattern inference theory is that Bayesian decision theory provides the best language both to describe complex patterns, and to model inferences about them. For us, the essence of a Bayesian view is not the emphasis on subjective prior probabilities, but rather that all variables are random variables. This assumption has ramifications for the central role, in perception, of generative (or synthetic) models of image patterns, as well as prior probability models of scene information. An emphasis on generative models, we believe, is essential because of the inherent complexity of the causal structure of high-dimensional image patterns. One must model how the multitude of variables (both the needed and unneeded variables for a task) interact to produce image data in order to understand how to decode those patterns. But perhaps equally

¹Our level of analysis falls between the computational/function and representation/algorithmic levels in the Marr hierarchy.

²Because it is rare to find a visual cue that is sufficiently reliable to unambiguously determine a perceived scene property, perception should be viewed as satisfying multiple constraints simultaneously. Examples are the constraint that light sources tend to be from above, or that a sharp image edge is more likely a reflectance or depth change than a shadow.

³These four classes are intended to apply generally to natural patterns of all sorts, and not just to visual patterns. For spatial vision, these classes would correspond to: blur and noise, geometric deformations, superposition (e.g. of basis images), and occlusions (Mumford, 1994).

importantly, the Bayesian view underscores the importance of *confidence-driven* visual processes. This latter idea leads us to the view that perception consists of sequences of computations on probabilities, rather than a series of estimations or decisions. We illustrate this with recent work on Bayes nets.

In the next section, we will show that pattern inference theory is a logical elaboration of ideal observer analysis in classical signal detection theory. However, it goes beyond standard applications of ideal observer analysis by emphasizing the need to take into account the full range of natural image patterns, and the intimate tie between perception and successful behavior.

2 Pattern inference theory: A generalization of ideal observers

Signal detection theory (SDT) was developed in the 1950's to model and analyze human sensory decisions given internal and external background noise (Peterson et al., 1954; Green and Swets, 1974). The theory combined earlier work in statistical decision theory (Neyman and Pearson, 1933; Wald, 1939, 1950; Grenander, 1950) with communication systems theory (Shannon and Weaver, 1949; Rice, 1944). Signal detection theory made two fundamental contributions to our understanding of human perception. First, statistical decision theory showed how to analyze the internal processing of sensory decisions. The application of statistical decision theory to psychophysics showed that sensory decisions were determined by two experimentally separable factors: sensitivity (related to an inferred internal signal-to-noise ratio) and the decision criterion. Second, communication theory showed that there were inherent physical limits to the reliability of information transmission, and thus detection, independent of the specific implementation of the detector, i.e. whether it be physical or biological. These limits can be modeled by a mathematically defined ideal observer, which provides a *quantitative* computational theory for the information in a task. For the ideal observer, the signal-to-noise ratio can be obtained from direct measurements of the variations in the transmitted signal. The ideal observer presaged Marr's ideas of a computational theory for an information processing task, as distinct from the algorithm and implementation to carry it out (Marr, 1982). The top panel of figure (1) illustrates the basic causal structure for the "signal plus noise" problem in classical signal detection theory.

Experimental studies of human perceptual behavior are often left with a crucial, but unanswered question: To what extent is the measured performance limited by the information in the task rather than by the perceptual system itself? Answers to this question are critical for understanding the relationship between perceptual behavior and its underlying biological mechanisms. Signal detection theory provided an answer through ideal observer analysis. One of the first applications of the ideal observer in vision was the determination of the quantum efficiency of human light discrimination (Barlow, 1962). By considering both the external and internal sources of variability, Barlow showed that an ideal photon detector could get by with about one tenth the number of photons as a human for the same combination of hit and correct rejection rates. This success of classical signal detection theory demonstrated the need for probability in theories of visual performance, because light transmission is fundamentally stochastic (emission and absorption are Poisson processes) and any real light measurement device introduces further noise.

The example of ideal observer analysis of light detection further illustrates a fundamental strategy for studying perception, consisting of three modeling domains. First, how does the signal (i.e. light switch set to "bright" or "dim") get encoded into intensity changes in the image? The answer must deal with light variations due to quantal fluctuations. Second, how should the received image data be decoded to

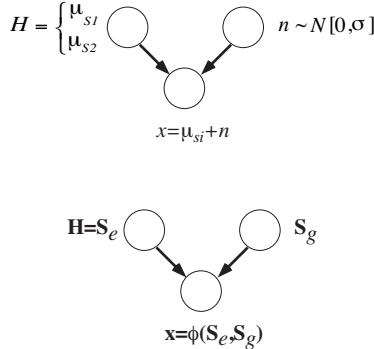


Figure 1: The top panel shows an example of generative graph structure for an ideal observer problem in classical signal detection theory (SDT). The data are determined by the signal hypotheses plus (usually additive gaussian) noise. Knowledge is represented by the joint probability $p(x, u, n)$. The lower panel shows a simplified example of the generative structure for perceptual inference from a pattern inference theory perspective. The image measurements (x) are determined by a typically non-linear function (ϕ) of primary signal variables (S_e) and confounding secondary variables (S_g). Knowledge is represented by the joint probability $p(x, S_e, S_g)$. Both scene and image variables can be high dimensional vectors. In general, the causal structure of natural image patterns is more complex and consequently requires elaboration of its graphical representation (see Section 3.4). For SDT and pattern inference theory, the task is to make a decision about the signal hypotheses or primary signal variables, while discounting the noise or secondary variables. Thus optimal perceptual decisions are determined by $p(x, S_e)$, which is derived by summing over the secondary variables (i.e. marginalizing with respect to the secondary variables): $\int_{S_g} p(x, S_e, S_g) dS_g$.

do the best job at inferring which signal was transmitted? Answers to this question rely on theories of ideal observers, or more generally of optimal inference. Third, how does one compare human and ideal performance? This requires common performance measures on the same task.

2.0.1 Limitations of Signal Detection Theory for the Grand Challenge

Despite its successes, signal detection theory as typically applied in vision falls short when faced with our grand challenge. Define perceptual signals to be some underlying causes of image data that are required for a visual behavior. These signals include the shapes, positions, and material of objects. The first problem is that natural perceptual signals are complex, high-dimensional functions of image intensities. In typical applications of SDT and classical ideal observer analysis to visual psychophysics, the input data, the noise, and the signal, are treated as the same “stuff”. For example, in contrast detection, the input data is signal plus noise (Kersten, 1984). The signal is based on a physical quantity (luminance) as a function of time and/or space), the noise is either physical contrast fluctuations, or internal variability treated as equivalent to the physical noise (Pelli, 1990). Perceptual decisions are typically limited to information which is explicit in the decoded signal. So to answer the question, Does the signal image have more light intensity than another?, the decoder simply measures whether the image intensity is bigger.

We need a theoretical framework for which the signals can be any properties of the world useful for the visual behavior; for example, estimates of object shape and surface motion are crucial for actions such as recognition and navigation, but they are not simple functions of light intensity. Natural images are high-dimensional functions of useful signals, and arriving at decoding functions relating image measurements

to these signals is a major theoretical challenge. Both of these problems are expressible in terms of pattern inference theory.

In signal detection theory, the non-signal causes of the input pattern are called noise. A second problem, related to the first, is that “noise” in the perception of natural images is not simple. Useful information is confounded by more than added external or internal image intensity noise. Uncertainty is due to both variations in unneeded scene variables as well as by the fact that multiple scene descriptions can produce the same image data. In contrast to the above example of contrast detection, consider the problem of 3D shape discrimination in everyday vision. The signal is shape, but the counterpart to the noise is very different stuff, and includes variation in viewpoint, illumination, other occluding objects, and material (Liu et al., 1995). Further, although the discrimination decision may be able to rely on a primary image measurement that is explicit in the image (e.g. a contour description), this is rare. Because of projection and the confounding variables, the true 3D shape is not explicit in any simple image measurement.

Pattern inference theory deals directly with the problem of multiple and diverse causes of image variation by modeling the generative process of image formation. Below, we distinguish between the needed *primary* and unneeded *secondary* variables.⁴ The primary variables are those which the system’s function is designed to estimate. By contrast, the secondary variables are not estimated but neither are they ignored, and there are principled methods for getting rid of unwanted variables. It should be emphasized that the distinction between primary and secondary depends on the specific task the system is designed to solve. Variables which are secondary for one task may be primary for another. For example, estimating the illumination is unimportant for many visual tasks and so illumination variables are treated as secondary. The theory of generic views treats viewpoint as a secondary variable, enabling resolution of ambiguities in shape perception (Nakayama and Shimojo, 1992; Freeman, 1994). Light direction as a secondary variable can be used to obtain a unique estimate of depth from cast shadows (Kersten, 1999). There is a close connection between the task (discussed in Section 4 below) and the statistical structure of the estimation problem (Schrater and Kersten, 2000).

A third limitation is that natural images are not linear combinations of their signals, and that the probabilities describing the signal and image variables are not gaussian. Much of the success of signal detection theory has rested on an assumption of linearity: the input is the sum of the signal and the noise. Except in rare instances (e.g. contrast detection limited by photon fluctuations at high light levels), natural perceptual tasks involve inputs which are non-linear functions of the signals and the noise (or secondary variables). For example, light intensity is a non-linear function of object shape, reflectance, and illumination.

There is a close relationship between linearity and the assumption that the random variables of interest are Gaussian⁵. Although classical signal detection explored the implications of non-Gaussian processes (Egan, 1975), most applications of signal detection theory to vision have typically approximated noise variations as Gaussian processes. A Gaussian approximation works very well in certain domains (as an approximation to Poisson light emission), but is extremely limited as a model of scene variability. Both the linear and Gaussian assumptions have had a striking success in the general problem of modeling human perceptual and cognitive decisions, where the variability is inside the observer (Green and Swets, 1974; Swets, 1988). But the Gaussian assumption generally fails when modeling external variability. For example, whenever a probability density involves more than second-order correlations,

⁴Primary and secondary variables have also been referred to as explicit and generic (or nuisance) variables, respectively.

⁵Because the log of a multi-variate Gaussian is quadratic, extrema can be found using linear estimators.

a multi-variate Gaussian model is no longer adequate. Image samples from Gaussian models of natural images fail to capture the rich structure of natural textures (Knill et al., 1990). Simple image measurements, such as those made by simple cells of the visual cortex are highly non-Gaussian (Field, 1987). A goal of pattern inference theory is to let the vision problem determine the distributions.

Fourthly, perception involves more tasks than classification. Not surprisingly, for signal detection theory, the primary focus is on signal detection—was the signal sent or not? Perception involves a larger class of tasks: classification at several levels of abstraction, estimation, learning, and control. Past applications of signal detection theory have successively handled certain kinds of abstraction (e.g. “is any one of 100 signals there or not?” or “which of 100 known signals was sent?”) as well as estimation (Van Trees, 1968); but we also require a framework that can handle diverse tasks from continuous estimations (e.g. of distance, shape, and their associations) to more complex categorical decisions: e.g. is the input pattern due to a cat, a dog, or “my cat”? Tools for the former build on classical estimation theory, but include recent work on hidden markov models. The latter requires additional tools, such as flexible template theories to model shape abstraction. A mathematical framework for perception requires tools for the generalization of ideal observers for the functional complex tasks of natural perception. Defining primary and secondary variables is part of task specification, and pattern inference theory handles this by incorporating decision theory to define a risk function (Section 4).

Finally, we note that most of the interesting perceptual knowledge on priors and utility is implicit. Signal detection theory grew out of earlier work on decision theory. Two important components of decision theory are the specification of prior probabilities of scene properties or signals and the costs and benefits of actions, through a risk or cost function. In most applications of SDT, it has been the experimenter that manipulates the priors and the cost functions. The human observer is often aware of the changes, and can adopt a conscious strategy to take these into account. We argue that the most important perceptual priors are largely determined by the structure of the environment and can, in principle, be modeled independently of perceptual inference (i.e. in the synthesis phase of study)⁶. Modeling priors (e.g. through density estimation) is a hard theoretical problem in and of itself, especially because of the large number of potential interactions. In classical SDT, probabilities are typically specified over small dimensional spaces. The costs and benefits are inherent to the type of perceptual task, and determine the primary and secondary variables. Thus, to elaborate on Helmholtz’s definition of perception: *perception is (largely) unconscious inference involving unconscious priors, and unconscious cost functions.*

Thanks to the successes of signal detection theory, we know that perception is limited by two factors: 1) the available information for reliable learning, inference, and action; 2) brain mechanisms to process that information. But one of the principal differences between classical SDT and pattern inference theory is the greater emphasis on modeling the external limits to inference, including both synthesis and optimal decoding. Both problems are clearly challenging, and computer vision has shown that the second problem is surprisingly hard. We agree with Marr when he wrote in 1982: “...the nature of the computations that underlie perception depends more upon the computational problems that have to be solved than upon the particular hardware in which their solutions are implemented.” Theories of human perceptual inference require an understanding of the limits of perceptual inference through optimal decoding theories (Barlow, 1981; Geisler, 1989). These theories, in turn, require an understanding of the transformations and variations introduced in pattern formation. We will argue here that the

⁶We emphasize an empirical Bayesian approach in which, as is discussed in Section 6, one can test an hypothesis relating a subjective prior to an objective prior.

structure of the visual information for function is best modeled in terms of its probabilistic structure, and that as a consequence any successful system must reflect the constraints in that structure, and further that its computations should be in terms of probability operations.

So, in the next section, we focus on the first problem: How can we model the information required for a task? This modeling problem can be broken down into: a) synthesis, modeling the structure of pattern information in natural images; and b) analysis, modeling the task and extracting useful pattern structures.

3 Encoding of scenes in images: Modeling image pattern synthesis

Computer vision has emphasized the difficulty of image understanding, which involves decoding images to find the scene variables causing the image measurements. Although, a great deal of progress has been made in computer vision, the best systems are typically quite constrained in their domain of applicability (e.g. letter recognition, tracking, structure from rigid body motions, etc.). The focus has understandably been on decoding—e.g. solving the inverse optics problem. However, the success of image decoding depends crucially on understanding the encoding. Although the computational challenge of image understanding is widely appreciated, the difficulty and issues of image pattern synthesis are less so.

How do we model the information images contain about scene properties? Following Shannon (1949), the answer is through probability distributions. Treating perception as a communication problem, we identify certain scene variables S as the messages, and the image formation and measurement mapping as the channel $p(I|S)$, by which we receive the encoded messages I . Given this identification, we can use information theoretic ideas to quantify the information that I gives about S as the transinformation

$$I(S; I) = H(I) - H(I|S) = E_{p(I)}[-\log p(I)] - E_{p(S,I)}[-\log p(I|S)].$$

These entropies are determined by $p(I) = \int_S p(S)p(I|S)dS$, the likelihood $p(I|S)$, and the prior $p(S)$.⁷ Thus, the physics of materials, optics, light, and image measurement, which determine the likelihood, just scratch the surface of what is required to model image encoding. In addition, we need to understand the types of patterns and transformations that result from the fact that images are caused by a structured world of events and potentialities for an agent, which is captured in $p(S)$. While probability and information theory provide the tools for understanding image encoding, constructing theories with these tools requires work. Let's look at the framework, tools, and principles for theory construction.

⁷For simplicity, we've restricted our expressions to probability densities on continuous random, rather than discrete, random variables. There are well-known subtleties in translating results between discrete probabilities and continuous densities. Examples: 1) A change of representation (e.g. changing distance to vergence angle) will in general change the form of the density—e.g. change a uniform density into a non-uniform one. 2) Entropy for continuous variables is inherently relative, and thus transinformation is more useful (Cover and Joy, 1991). 3) If the range of a random variable is unknown, then the principle of insufficient reason leads to “improper” priors (Berger, 1985).

3.1 Essence of Bayes: Everything is a random variable

A key starting assumption is that all variables are random variables, and that the knowledge required for visual function is specified by the joint probability $p(S_e, S_g, I)$. The basic ingredients are variable classes: image measurement data (I), variables specifying relevant scene attributes (S_e), and the confounding variables (S_g). All these variables are random variables, and thus subject to the laws of the probability calculus, including Bayes theorem. So for pattern inference theory, a Bayesian view is more than acknowledging the role of priors, but also emphasizes the redundancy structure of images, and the importance of the generative process of visual pattern formation, expressible as a graphical model. Thus the essence of a Bayesian theory of perception is more than applying Bayes' rule to infer scene properties from images, or that likelihoods are tweaked by prior and labile subjective "biases". This interpretation (*Myth 1: Bayesian models of perception are distinct only by virtue of emphasis on modeling priors*⁸) would miss the point of our view of pattern inference theory approach to perception. By starting with a model space completely determined by the joint probability, $p(S_e, S_g, I)$, we have the foundation to understand:

1) input image redundancy, through:

$$p(I) = \int p(S_e, S_g, I) dS_e dS_g$$

2) scene structure, through:

$$p(S_e, S_g) = \int p(S_e, S_g, I) dI$$

and

3) inference, through:

$$p(S_e, I) = \int p(S_e, S_g, I) dS_g$$

Of course, modeling $p(S_e, S_g, I)$ in general may pose an insurmountable challenge. But there is reason for optimism, and recent work in density estimation and image statistics suggest that tractable high-dimensional models may be possible (Zhu et al., 1997; Zhu and Mumford, 1998; Zhu, 1999; Simoncelli, 1997, 1998).

The key point is that necessary knowledge to characterize the perceptual problem is specified by a joint probability over the given data (usually image measurements, but could include contextual conclusions drawn earlier or elsewhere), what the visual system needs (primary), and the variables that confound (secondary variables).

⁸At several points in this chapter, we address what we see as misconceptions of the Bayesian framework for vision. We identify these as "myths".

3.2 Basic operations on probabilities: Conditioning and Marginalizing

We really have only two basic computations on probabilities, which follow from the two basic rules of probability—the sum and product rules. Each of the rules has specific roles in an inference computation, related to the kind of variable in the inference. When inferring the values of a set of variables S_e , the remaining variables come in two types: those which we don't know and don't care to know, and those which are known either by sensory measurement or a priori. How is the joint probability affected by this knowledge? The answer is to sum over the unneeded variables (marginalization), and divide the joint by the probability of the known ones (conditioning).

1) Marginalization: Presuming the utility of only a subset of the scene variables (which we treat in Section 4), the values of some variables S_g are not known, and we don't care to know them. Marginalization is the proper way to remove the effect of these secondary, unknown, unwanted variables:

$$P(S_e, I) = \int_{S_g} P(S_e, S_g, I) dS_g$$

The reason we marginalize is that being unneeded doesn't mean these variables should be ignored! Most of the time, the unwanted variables (e.g. viewpoint) crucially contribute to the generation of the possible images, and hence cannot be ignored. The marginalization approach contrasts with traditional modular studies of vision, in which most of the unneeded variables for a given module are left out of the discussion entirely (e.g. independent estimation of reflectance, shape, and illumination). Often, the modularity is adopted based on general practical and theoretical arguments. Our position is not that we forgo modularity, but rather that modularity be grounded in the statistical structure of the problem, rather than by what the theorist finds convenient (Schrater and Kersten, 2000). It is important to emphasize that this approach does not necessitate that marginalizations are executed on-line by the brain. The effects of marginalization could be built directly into the inference algorithm avoiding the need for perception to have an explicit representation of the unneeded variables.

2) Conditioning: Some of our variables are known, through data measurements, or a priori assumptions. In either case, once we know something about the variables, we base our inferences on this knowledge by conditioning the joint distribution on the known information:

$$P(S|I) = P(S, I)/P(I)$$

The way Bayes' rule comes into the picture is that it is often easier to separately model image formation and the prior model for the causal factors. Bayes' rule is a straightforward application of the product rule to $P(S, I) = P(I|S)P(S)$:

$$P(S|I) = P(I|S)P(S)/P(I)$$

The likelihood $P(I|S)$ is determined by the generative image formation model which produces image measurements from a scene description. The generative model produces the image patterns, and consists of the scene prior, and the image formation model. The likelihood is easier to model because we are conditioning on the scene, and the image is a well-defined function of the scene—forward optics plus measurement noise. Although the likelihood and prior terms are logically separable, the division has

little bearing on the algorithmic implementation. When it comes to inference, Bayes is neutral with respect to whether a priori knowledge is used in a bottom-up or top-down fashion (*Myth 2: Priors are top-down.*). The regularizers in computer vision can be expressed as priors, and these are typically instantiated as bottom-up constraints (e.g. weights in feedforward networks (Poggio et al., 1985; Poggio and Girosi, 1990)).

Why should scene variables be treated probabilistically? In contrast to subjective Bayesian applications (*Myth 3: Priors only refer to subjective, and perhaps conscious biases.*), prior probabilities on scene variables are objectively quantifiable. They result from physical processes (e.g. material properties such as albedo and plasticity covary due to common dependence on the substance, such as metal), and from the relative frequencies of the scene variables in the observer’s environment. Thus, vision modelers have a big advantage over stock market analysts: they have a better idea of what the functionally important scene causes are, and can hypothesize and test probability density models of scene variables, independent of visual inference. They can also test the extent to which vision respects the constraints in the prior model (see Section 6). Why is probability essential for modeling pattern synthesis? Because an infinite set of scenes can produce a given image. Thus, in the decoding problem it is essential to have a model of the generative structure of images, given by $p(S|I)$ and $p(S)$. Below we discuss how several kinds of generative processes produce characteristic image patterns.

3.3 Generative models in vision

Functional vision depends on the kind of abstraction required for the task at hand. But psychological abstractions such as scene categories, object concepts, and affordances rest on the existence of objective world structure. Without such structure, there would be no support for reliable inferences—in fact, there would be no basis for consistent action in a world in which each image is independent of any previous ones. From this perspective, it is not unreasonable for an otherwise functional visual system to hallucinate in response to visual noise, because the best world interpretations will be structured. Thus, understanding the objective generative structure is necessary although not sufficient for an account of human visual perception⁹. However, a central theme of this chapter is the importance of understanding the objective generative processes of the images received. It is an intriguing scientific question as to the degree with which perceptual inference mechanisms mirror or recapitulate the generative image structure. Theories of back-projections in visual cortex rest on internal generative processes to deal with “explaining away” (Dayan et al., 1995; Hinton and Ghahramani, 1997) (see Section 4.3), the related idea of model validation through residual calculation (Mumford, 1992, 1994), and predictive coding (Rao and Ballard, 1999). As we discuss later, the task itself refines our model of the relevant statistical structure through Bayesian modularity.

Visual perception deals with two broad classes of generative processes that produce *photometric* and *geometric* image variation. Further, it is useful to distinguish scene variations (knowledge in $p(S)$) from those of image formation (knowledge in $p(I|S)$). We postpone the discussion of the experimental implications of these variations until Section 6.

⁹This is one way of distinguishing the Bayesian perspective from a strict Gibsonian view which could be interpreted as assuming that objective structure is also sufficient to explain functional vision.



Figure 2: Illustrations of variations in scene variables. Top left: A collection of bicycles shows variations in geometry (size, shape) and albedo (paint patterns). Top right: Two images of the same bicycle from the same view but differing in articulation. Bottom left: A flat-tailed Gecko hides on a tree, showing how variation in skin pigment (albedo) can match the background pattern of the tree bark (Copyright Martin Kramer). Bottom Right: A river illustrates the complexities of spatial layout. The presence and directionality of the water is encoded in the complex array of specularities and light scatter, determined by the interaction between light source, water surface fluctuation, and viewpoint.



Figure 3: Variations in illumination and viewing. Top Left: A young man's face shows shading variation due to the extrinsic shadows cast by leaves. Top Right: Reflection of a face on the glass door of a bookcase creates a transparent image. Bottom Left: Two images of the same bicycle differing in viewpoint. Bottom right: A deer hides behind foliage, illustrating occlusion/background clutter (Copyright Mark Brady).

3.3.1 Object and scene variations

A logical prerequisite for a full understanding of image variation is a study of the nature of illumination, surface reflectivities, object geometry, and scene structure quite independently of the properties of the sense organs. Consider geometrical object variations that occur for a single object. An individual object, such as a pair of scissors, or a particular human body consists of parts with a range of possible articulations. The modeling problem is of significant interest in computer graphic synthesis because it provides the means to model the transformations, and ultimately characterize the probabilities of particular articulations and actions.

Objects (and scenes) can be categorized at more abstract levels, such as “dogs” or “books” (Bobick and Richards, 1986; Rosch et al., 1976). Examples of sources of *within-class* scatter include geometric variations that occur between different members of the same species, vehicles, or computer keyboards. Certain types of within-class geometric variation (e.g. “cats”) can be modeled in terms of prototypes together with a description of geometric deformations (Grenander et al., 1991; Yuille, 1991; Mumford, 1996) which admit a probability measure $p(S)$. For this sort of within-class variation, it may be possible to find $p(S)$ through probability density estimation on scene descriptions. Estimating prior densities (e.g. via Principal Components Analysis or PCA) for the distribution of facial surfaces (variations across human face shapes) is now possible due to advances in technology for measuring depth maps (Atick et al., 1996; Vetter and Troje, 1997). Material or albedo variation also occurs across an object set—e.g. the set of books, with different covers. And of course there are mixtures of geometrical and photometrical effects, such as within-species variation among dogs. There is a considerable body of work on biological morphometrics whose goal is to understand the transformations connecting objects within groups (Bookstein, 1978, 1997; Kendall, 1989). Origin of concepts at certain levels may lie in the generative structure of objects, and debate has occurred as to whether an entry-level object concept is based on a prototype with (possibly) a metric model of variation, or a description of the structural relationships between parts. We touch on this point later in the context of object recognition models.

“Schemas” are an example of an even higher level of organization involving *spatial layout*, which recognizes the spatial relationships between objects, and their contextual contingencies. The fact that perceptual judgments are strongly influenced by scene context, (e.g. forest, office, or grocery store scene), suggests that $p(S)$ is not at all uniform across spatial layout, but rather is highly ‘spiked’ which allows scene type recognition and its exploitation for scene analysis. See figure 2 for examples of scene variable variations.

3.3.2 Effects in the image

At the most proximal stage, the images projected into the eyes are transformed by the optics, sampled by the retina, and have noise added to them. These operations produce the well-studied photometric variations of *luminance noise* and *blurring* in the images..

Due to the additivity of light and the approximate linearity of reflection, photometric variations due to *illumination* change are approximately linear so that under fixed view, an arbitrary lighting condition can be approximated by the weighted sum of relatively few basis images (Epstein et al., 1995). Further, it has been shown that the images of an object fall on or near a cone in image space (Belhumeur and Kriegman, 1996). Cast shadows are another form of illumination variation resulting from the occlusion of a light source from a surface. Specularity in an image is an interaction between material, shape, and

viewpoint. *Surface transparency* is another source of photometric variation. Its effect in the image can be either additive or multiplicative (Kersten, 1991). One form of additive transparency results from the combination of reflections in a store-front window.

Variations in observer viewpoint (i.e. viewing distance and direction) produce *geometric deformations* in the image. The utility of multiple-scale analysis in human and machine vision is in part a consequence of the distribution of translations in depth of the eye (or camera) viewpoint. Over small changes in viewing variables, the image variations are fairly smooth, although rotations around the viewing sphere can cause large changes in the images due to significant self-occlusions. If we were only concerned with geometry, viewpoint variations and variation in object position and orientation would produce the same set of images. However, illumination interacts with viewpoint so that view rotation is only equivalent to object rotation if the lighting is rigidly attached to the viewer's frame of reference. Rotations in depth cause particularly challenging image variations for object recognition that we briefly discuss later.

The distribution of multiple objects in a scene affects the images of objects through *occlusion* and clutter. Because of the nature of imaging, the local correlations in surface features typically carry over to local image features. However, occlusion of one object by another breaks up the image into disconnected patches. Further, patches widely separated in the image can be statistically related, and the challenge is to link the appropriate image measurements likely to belong to the same objects. Like occlusion, *background* is a significant confounding source of image variation that thwarts segmentation. The intensity edges at the bounding contours of an object can vary substantially as the background is changed, even if the view and lighting remain the same. See figure 3 for examples of illumination and viewing effects on images.

Occlusion is the result of the distribution of the kinds and spatial arrangements of objects within a scene relative to the viewpoint. But the *spatial layouts* of schemas also generate statistical dependence in images. Temporal variation in images is induced by object motion and observer actions (Bobick, 1997). Thus the spatio-temporal image distribution is affected by the distribution of observer *actions*, and object *dynamics* (e.g. freeway driving).

3.4 Graphical models of statistical structure

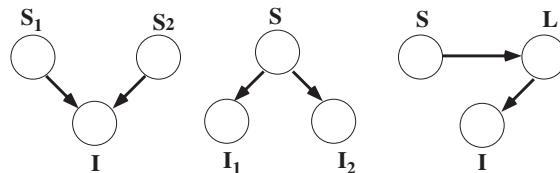


Figure 4: Components of the generative structure for image patterns involve converging, diverging, and intermediate nodes. For example, these could correspond to: multiple (scene) causes $\{S_1, S_2\}$ giving rise to the same image measurement, I ; one cause, S influencing more than one image measurement, $\{I_1, I_2\}$; a scene (or other) cause S , influencing an image measurement through an intermediate variable L .

In general, natural image pattern formation is specified by a high-dimensional joint probability, requiring an elaboration of the causal structure that is more complex than the simplified model in the bottom panel of figure 1. The idea is to represent the probabilistic structure of the joint distribution $P(S, I)$ by a Bayes net (Pearl, 1988; Ripley, 1996), which is simply a graphical model that expresses how variables influence each other. There are just three basic building blocks: converging, diverging, and

intermediate nodes. For example, multiple (e.g. scene) variables causing a given image measurement, a single variable producing multiple image measurements, or a cause indirectly influencing an image measurement through an intermediate variable (see figure 4). These types of influence provide a first step towards modeling the joint distribution and, as we describe in Section 4 below, the means to efficiently compute probabilities of the unknown variables given known values.

Influences between variables are represented by conditioning, and a graphical model expresses the conditional independencies between variables. Two random variables may only become independent, however, once the value of some third variable is known. This is called *conditional independence*.¹⁰

Using labels to represent variables and arrows to represent conditioning (with $a \rightarrow b$ indicating b is conditioned on a ¹¹), independence can be represented by the absence of connections between variables. For example, if the joint probability $p(a, b, c, d, e, f, g)$ factors by independence into $p(a, b, c, d, e, f, g) = p(a)p(b)p(c|d)p(d)p(e|a, b)p(f|b, c)p(g|d)$, then the variables can be represented by the graph in figure 5. Had the variables factored into two independent groups the graph would have shown two separate nets. The example graph can represent a Bayes network for computing structure from stereo and texture if we allow some of the nodes to represent multiple variables. To illustrate, let the node a represent the geometric and material causes of a particular image texture, and e represent the collection of texture measurements made by the observer. The node b represents absolute depth from the observer and is the variable of interest for the task. The horizontal and vertical disparity measurements are bundled into f , which depends on both the depth variable b and the direction and distance of the observer's fixation point, c , in space. The fixation point distance is determined by the convergence angle, d , between the eyes. The convergence angle can be inferred from non-visual proprioceptive feedback from the eyes represented by the data variable g .

Note that the graphical structure captures the structure of the data formation. The top layer of the graph represents the scene and viewing variables, whose causal effect on the sensory data in the bottom layer is represented by the directed arrows.

4 Optimal decoding: Modeling the tasks

The basic tenet is that perception enables successful behavior, and thus any decoding scheme is designed to extract useful information about the true state of the world. But the essence of decision theory analysis is the trade-off between truth and utility. A complete characterization of optimal behavior cannot dispense with either dimension. Even the simple problem of deciding whether a flash of light is bright or dim is only a useful visual function, if the task is to decide whether one or the other determines a true state of the world. Was the light switch set to high or low? Was the object closer or nearer? The fundamental computational problem of vision is: given visual data, how can the system determine the environmental causes of that data, when it is confounded by other variables. If one accepts this, then we can make the case that visual perception is fundamentally image decoding. But whether to draw an inference, or the precision with which it must be drawn is determined by the visual function. As

¹⁰Two random variables are independent if and only if their joint probability is equal to the product of their individual probabilities. Thus, if $p(A, B) = p(A)p(B)$, then A and B are independent. If $p(A, B|C) = p(A|C)p(B|C)$, then A and B are conditionally independent. When corn prices drop in the summer, hay fever incidence goes up. However, if the joint on corn price and hay fever is conditioned on “ideal weather for corn and ragweed”, the correlation between corn prices and hay fever drops. Corn price and hay fever symptoms are conditionally independent.

¹¹In graph theory, a is called the *parent* of b

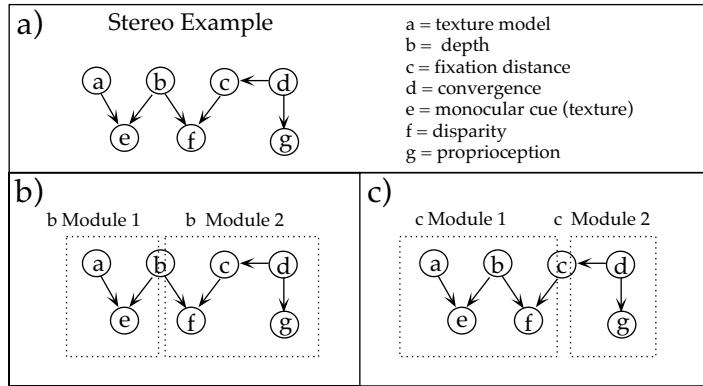


Figure 5: Example Bayes Net. **a)** Bayes net representing the factored distribution $p(a, b, c, d, e, f, g) = p(a)p(b)p(c|d)p(d)p(e|a, b)p(f|b, c)p(g|d)$. The graphical model can express the probabilistic structure of depth from stereo and texture inference. **b)** When estimating the depth variable b , the Net can be decomposed into two separate depth modules (depth from texture and depth from stereo). The dashed boxes show the modules. **c)** When estimating the fixation distance c , the net can be decomposed into two separate distance modules (distance from texture and stereo, and distance from proprioception). Note that the left hand side of the graph does not decompose as before. This illustrates *Bayesian modularity* (see Section 4.3).

with any decoding system, perception operates with target goals that depend on the task. A complete theory of vision needs to account for three classes of behavioral tasks:

- 1) The visual system draws discrete (categorical) conclusions about objective world. These decisions invariably involve taking into account potentially large variations in confounding secondary variables. For example, to reliably detect a face, a system must allow for variations in view, lighting, background, as well as the geometrical variations in individual facial shape, expression, and hair. Finer-grain identifications require more estimates of primary variables, and less marginalization with respect to secondary variables (Kersten, 1997). Because the causes are objective, decisions have a right or wrong answer. Further, the cost due to incorrect decisions can be large or small. A mistake in animal identification can have serious consequences. Failing to anticipate the change in color of a sweater going from indoor to outdoor lighting may cause only mild social embarrassment, requiring little investment in perceptual (vs. learned cognitive) resources.
- 2) The visual system provides continuously valued estimations for actions. For example, visual information for depth and size determine the kinematics of reach and grasp. Like discrete decisions, estimations can have degrees of utility.
- 3) The visual system adapts to environmental contingencies in the images received. This adaptation is at longer time scales than inference required for perceptual problem solving, occurring over both phylogenetic and ontogenetic scales. One form of adaptation requires implicit probability density estimation.

Can we describe these processes from the point of view of pattern inference theory—i.e. as image decoding by means of probability computations? To do so requires a probabilistic model of tasks. We consider a task as specifying four things, the required or primary set of scene variables S_e , the nuisance or secondary scene variables S_g , the scene variables which are presumed known S_f , and the decision to be made. Each of the four components of a task plays a role in determining the structure of the optimal inference computation. First, we review how to model the decision as a risk functional on the posterior distribution, then we show that S_e and S_f can be used to simplify the joint distribution through independence relations, while S_g and the decision rule can make one choice of S_e simpler than another.

Bayesian decision theory provides a precise language to model the costs of errors determined by the choice of visual task (Yuille and Blthoff, 1996; Brainard and Freeman, 1997). The cost or *risk* $R(\Sigma; I)$ of guessing Σ when the image measurement is I is defined as the expected *loss*:

$$R(\Sigma; I) = \int_S L(\Sigma, S) P(S | I) dS,$$

with respect to the posterior probability, $P(S|I)$. The best interpretation of the image can then be made by finding the Σ which minimizes the risk function. The loss function $L(\Sigma, S)$ specifies the cost of guessing Σ when the scene variable is S . One possible loss function is $-\delta(\Sigma - S)$. In this case the risk becomes $R(\Sigma; I) = -P(\Sigma | I)$, and then the best strategy is to pick the most likely interpretation. This is standard *maximum a posteriori estimation* (MAP). A second kind of loss function assumes that costs are constant over all guesses of a variable. This is equivalent to marginalization of the posterior with respect to that variable.

The introduction of a cost function makes Bayesian decision theory an extremely general theoretical tool. However, this flexibility has drawbacks from a scientific perspective. We could potentially introduce a loss function for each scene variable, which makes it impractical to independently test cost function hypotheses empirically—and we are stuck with an additional set of free parameters. However, we can

achieve modeling economy by assuming the delta function or constant loss functions depending on whether the variable is needed (primary) or not. Thus, we advocate initially constructing simpler Bayesian theories in which we estimate the most probable relevant scene value (MAP estimation), while marginalizing with respect to the irrelevant generic variables. Bloj and colleagues have a recent example of this strategy applied to the interaction of color and shape (Bloj et al., 1999).

We now describe how the statistical structure and task interact in determining the inference computations. While the statistical structure of the joint distribution determines which variables interact, the choice of decision rule and marginalization variables determine the details of how they interact. In the next section, we show how the task, in choosing the relevant variables, partitions the scene variables through statistical independence.

4.1 Partitioning the scene causes: Task dependency and conditional independence

Considering a single task allows us to focus our attention on a particular set of variables S_e . In some cases, we may be justified in ignoring a number of scene properties irrelevant to the task. This idea can be expressed in terms of the distributions through statistical independence. We may factor $p(S, I)$ into two parts, one of which contains all the variables which are statistically independent of S_e and the other which contains all of the dependent variables, $p(S, I) = p(I_{ind}|S_{ind})p(I_{dep}|S_{dep})p(S_{ind})p(S_{dep})$ ¹². In terms of a graphical model, this partitioning corresponds to unconnected sub-graphs. Specifying a task restricts our base of inference to $p(I_{dep}, S_{dep})$.

In addition, the nature of a task or context fixes some of the scene variables S_f . For instance, if an observer is doing quality checking on an assembly line, then the lighting variables and viewpoint can be considered fixed. Note that constraints used to regularize vision problems can often be expressed as fixing a set of scene variables. For instance, in a world of polynomial surfaces, the constraint that the task only involves flat surfaces can be rephrased as all non-linear polynomial coefficients are fixed at zero.

Since the variables in S_f are presumed known, we can subdivide the dependent variables still further, $S_{dep} \rightarrow S'_{dep}, S_f$ and condition $p(I_{dep}, S_{dep})$ on S_f , $p(I_{dep}, S'_{dep}|S_f)$, which increases the statistical independence of the variables. This is true because variables which are not statistically independent, because they are dependent on a common variable, become independent when conditioned on the common variable. Thus we expect the conditional distribution to further decompose into relevant and irrelevant scene variable components.

Thus given the task, we can first factor $p(S, I|S_f) = \prod_{i=1}^N p(S_i, I|S_f)$. To do inference we need only consider the factors in which the S_i contain the variables in S_e . Let S_j denote the minimal set of statistically dependent variables containing S_e . The variables in S_j excluding S_e are just the secondary variables S_g . Then, $p(S_e, S_g, I|S_f)$ contains all the information we need to perform the inference task, and has automatically specified the task relevant and irrelevant variables, i.e. the primary and secondary variables. Thus the independence structure determines which variables should be involved in an inference computation. This is an important issue for modeling cue integration.

In terms of graphical models, the set of variables S_e and S_g for the task have the property that they are connected by the image data. In other words, S_e and S_g are both involved in generating the image data. The basic generative structure of the perceptual inference problem is illustrated in the lower panel in

¹²For notational convenience, here we use S to indicate the set of scene variables to be partitioned, $\{S\}$.

	Object perception		Spatial layout		
	Object-centered (object recognition)		World-centered	Observer-centered (hand action)	
	Basic-level	Subordinate-level		Planning	Reach
Shape	E	E	G	G	G
Material	G	E	G	G	G
Articulation	G	E	G	G	E
Viewpoint	G	G	G	E	G
Relative position	G	G	E	G	G
Illumination	G	G	G	G	G

Table 1: Table illustrating how the visual task partitions the scene variables into primary (E) and secondary (G) variables. The pattern of image intensities is determined by all of the scene variables, object shape, material, object articulation (e.g. body limb movements or facial expression), viewpoint, relative position between objects, and illumination. Basic-level recognition involves more abstract categorization (e.g. dog vs. cat) than subordinate-level recognition (Doberman vs. Dachshund), and is typically thought to be shape-based, with material properties such as fur color discounted. Finer-grain subordinate-level recognition requires estimates of shape and material.

figure 1 from the point of view of pattern inference theory. Comparing this diagram to the generative diagram for the standard signal detection theory above it, we can better see how pattern inference theory is a generalization of the typical way of using signal detection theory. In most applications of SDT to vision, the image data are generated by signals plus noise, which allow us to identify S_e as the signal set, and S_g as the noise. Thus, one of the key ideas of pattern inference theory is that unwanted variables act like noise in the context of a particular inference task. However, the noise is multivariate, highly structured and in general cannot be modeled by a unimodal distribution. While the set of generic variables, S_g , play the role of noise for one task, they form the “signal” for another task, because the distinction between primary and secondary depends on the visual function. What is a primary variable for one task may be secondary for another. Table 1 illustrates how various visual tasks determine the primary vs. secondary variables.

One of the consequences of deciding a task, is that ambiguity can be reduced through marginalization (Freeman, 1994; Knill et al., 1996b). The basic principle is: *perception’s model of the image measurement ((i.e. the generative consequence of the primary variable’s prediction of the image measurement) should be robust with respect to variations in the secondary variables.* In fact, the general viewpoint principle is a consequence of viewpoint being a secondary variable (Freeman, 1994).

4.2 Partitioning image measurements: Sufficient statistics

Once we have determined which scene variables are relevant to the task, the independence structure of $p(S_e, S_g, I|S_f)$ specifies the image measurements to make. Assuming we have a set of measurements $\{m_1(I), m_2(I), m_3(I), \dots\}$ which form a good code for $p(I)$, then we can determine which image measurements to use by partitioning the joint distribution. The joint distribution,

$$p(S_e, \{m_1(I), m_2(I), m_3(I), \dots\} | S_f)$$

will further factor into relevant and irrelevant image measurements, yielding a set M of measurements required for the task. If we inspect the posterior distribution needed for inference $p(S_e | M, S_f)$, we

can interpret the set M as the set of sufficient statistics for S_e , since $p(S_e|I, M, S_f) = p(S_e|M, S_f)$ fits the standard definition of a *sufficient statistic* (Duda and Hart, 1973). While many different sets of measurements can form sufficient statistics, *minimal sufficient statistics* are the smallest set of sufficient statistics and have the property that any other set of sufficient statistics are a function of them. This new perspective leads to the principle: a good image code for a visual system is one that forms a set of minimal sufficient statistics for the tasks the observer performs.

4.3 Putting the pieces together: Needed scene estimates, sufficient image measurements, and probability computation.

We have shown for optimal inference, how the choice of required variables determines which scene variables we need to consider through statistical independence, and the set of image measurements through the notion of sufficient statistics. We now illustrate how the variables interact in optimal inference, which is determined by the details of the generative model and the choice of loss functions. The generative model, in specifying how the secondary variables interact with the primary variables to produce an image, determines to a large extent how the primary and secondary variables interact in an inference computation. However, the choice of cost function, by specifying different costs for errors, modulates the relevance of errors induced by the ignorance of particular secondary variables.

To be more specific, we return to the generative model for texture, disparity and proprioceptive data. (figure 5), but now from the point of view of decoding—estimating depth from measurements of disparity and texture. In Bayesian inference, the change in certainty of the scene variables causing the image after receiving image data respects the generative model. Both prior knowledge and image measurements fix values in the network, and the problem is to update the probabilities of the remaining variables. Updating the probabilities is straightforward for simple networks, but requires more sophisticated techniques such as probability or belief propagation, or the junction-tree algorithm for more complex networks (Frey, 1998; Weiss, 1997; Jordan, 1998). The primary effect of receiving image data is to change the certainty of all the variables which could possibly generate the image data. One effect of having more than one image measurement is known as “explaining away” in Bayes nets. For example, suppose we observe that the texture measurements e are compressed in the y direction relative to an isotropic texture. The compression might be the result of our texture being non-isotropic (i.e. attributing the observation to the texture model a), it might be due to the surface having a depth gradient (i.e. attributing the measurement to the surface depth b), or it might be due to a little of both. Given only the texture measurement, the data supplies evidence for both a and b . However, if we have additional disparity data f which is consistent with a depth gradient, then our best inference is that both the texture compression and the disparity gradient are caused by a depth gradient. This second piece of information drives the additional inference that our texture model should be isotropic—a common depth gradient “explains away” the coincidence between the disparity gradient and the texture compression. Bayesian inference does this naturally by updating probabilities of each needed but unknown variable. The process of updating probabilities in a network is more powerful than estimating a single state. For example, if the random variables in the network are Gaussian, then updating probabilities requires new estimates of the mean *and* variance.

The task also affects the algorithmic structure. To illustrate, consider trying to do inference based on the total probability distribution. We would need to maintain a probability distribution on more than 7 dimensions (one for each node in the network plus the nodes with multiple variables). Thus, computing

using the entire distribution would be computationally prohibitive. However, the statistical independencies show a kind of modularity we call *Bayesian modularity*. In Bayesian modularity, the independence structure allows us to produce separate likelihood functions for the variable of interest, which can be combined by multiplication. For instance, if we are doing inference on b in the above example, $p(e|b) = \int_a p(e|a, b)da$ produces one likelihood function and $p(f, g|b) = \int_c [\int_d p(g|d)p(c|d)p(d)dd] p(f|b, c)dc$ produces the other. This division creates two 'modules' illustrated in figure 5b. The division also creates enormous computational savings, as we only need to maintain three likelihoods over two variables: $\{a, b\}$, $\{b, c\}$ & $\{c, d\}$. Modularity is modulated by the task. Figure 5c shows how Bayesian modularity changes as a function of which variables are estimated.

The quantitative influence of the data on the inference depends critically on both the likelihood and the knowledge we have about the secondary variables. The value of priors on secondary variables is clear, however the effect of likelihood is more subtle, as it depends on the number of possible scene causes for an image and the change in the image given a change in the scene variables. For example, Knill has shown that texture information is less reliable for frontal parallel surfaces than for strongly slanted surfaces because large changes in slant for fronto-parallel surfaces cause small changes in image texture compared to slant changes for strongly slanted surfaces (Knill, 1998b).

Now depending on our cost function, the two likelihood functions $p(e|b)$ and $p(f, g|b)$ for the depth b will have different influences on the decision. For example, consider a depth task in which the cost of depth errors is only high when the depth gradient is small (i.e. the surfaces are nearly fronto-parallel). In this case the depth from texture module will be nearly irrelevant to the decisions, because texture information is only reliable for large depth gradients (Knill, 1998b), whereas disparity information can be reliable for small depth gradients.

5 Learning generative structure

In pattern inference theory, learning is estimating the density $p(S, I)$, and discovering the appropriate cost function for the task. For example, learning to classify images of faces as male or female requires knowledge of intragender facial variability (i.e. $p(S)$), knowledge about how faces produce images (i.e. $p(I|S)$), and the decision boundary set by the cost of incorrectly identifying the faces. The two components, density estimation and cost function specification, have a rough correspondence to what we might call task-general and task-specific constraints respectively. Task-general constraints are those which hold irregardless of the specific nature of the task, which correspond to the fundamental constraints on inference set by the structure of the joint density. On the other hand, the choice of cost function is always task-specific, since it involves specifying the costs for a particular task. For generality, we focus on density estimation below.

It is one thing to talk about what one could do given the joint probability for a visual problem, and it is quite another matter to actually obtain it. High-dimensional probability density estimation is notoriously difficult. This observation has lead to radically different alternatives to learning, which place focus on the decision boundaries, largely ignoring the within-class structure (e.g. Support vector machines (Vapnik, 1995)). We discuss here several reasons to be optimistic.

An essential requirement for density estimation is to have a rich vocabulary of possible densities, which are typically parametric, from which a best fit to the image data can be achieved. The second requirement is having a sensible error metric to assess the best fitting density model. Zhu, Wu & Mumford

(1997) have developed a general method for density estimation based on the *Minimax Entropy Principle* which allows the consideration of both the best fitting model and what image measurements should be used. They assume that the density can be approximated well by a Gibbs distribution. Given a set of image measurements, they fit the best Gibbs distribution using the maximum entropy principle (Jaynes, 1957)¹³, which in essence chooses the least structured distribution consistent with the image measurements. They then use the Kullback-Leibler divergence to select between different models and sets of image measurements. Maximum entropy fits prevent model overfitting and choose the Gibbs distributions for which the set of image measurements are sufficient statistics.

Another approach to density estimation works by evaluating the *evidence* (MacKay, 1992). Let G represent an index across the set of generative models we are considering. Then we select the best fitting model by maximizing the evidence $p(G|I) = p(I|G)p(G)$, where $p(I|G) = \int_{S_G} p(I|S_G, G)p(S_G)dS_G$. Assuming we have a lot of image data, the prior across models does not matter much and the decision is based on $p(I|G)$. Choosing models by maximizing the evidence naturally instantiates Occam's Razor, i.e. models with lots of parameters are penalized (MacKay, 1992). Schwarz (Schwarz, 1978) has found an asymptotic approximation to $\log p(I|G)$ for well behaved priors which makes the penalty for the number of parameters of G explicit: $\log p(I|G) \simeq \log p(I|G, \hat{S}_G) - \frac{\log N}{2} \text{Dim}(G)$, where N is the number of training samples, \hat{S}_G is the maximum likelihood estimate of the scene parameters and $\text{Dim}(G)$ is the number of parameters for the model G . A similar formula arises from the *Minimum Description Length* (MDL) principle, which through Shannon's optimal coding theorem, is formally equivalent to MAP. While embodying Occam's Razor, evaluating the evidence works by choosing the model which is the best predictor of the data.

There have also been a few studies that try to directly learn a mapping from image measurements to scene descriptions (Freeman and Pasztor, 1999; Kersten et al., 1987). However, these approaches are limited in requiring the availability of sample pairs of scene and image data. While general methods could be used by the visual system for learning, the visual system may employ quite impoverished models of the joint density. The key point is that learning algorithms for both objective physical modeling or biological learning can be expressed in the Bayesian language of pattern inference theory.

6 Testing models of human perception

In order for the pattern inference theory approach to be useful, we need to be able to construct predictive theories of visual function which are amenable to experimental testing. While we have discussed the elements of constructing Bayesian theories throughout the paper, it is important to distinguish the role of the mathematical language from the elements of a theory of vision.

6.1 Pattern Inference Theories of Vision

How do Bayesian or pattern inference theories of vision differ from other theories (e.g. Gestalt)? The answer so far is that they express observer principles in terms of probabilities and cost functions. Thus, these theories will involve explicit statements about the scene variables and image measurements used,

¹³The Maximum Entropy Principle is a generalization of the symmetry principle in probability, and is also known as the principle of insufficient reason. For example, it says that one should assume a random variable is uniformly distributed over a known range unless there is sufficient reason to assume otherwise.

the prior probabilities on scene variables, the image formation and measurement model assumed by the observer, and the relative costs assigned to potential outcomes in a task. We also hope however, that pattern inference theory lead to a set of fundamental and deep principles akin to the laws of thermodynamics, also expressible in the same framework. The importance of such principles for scientific economy should not be underestimated. From the right first principles, an infinite set of experimentally testable consequences can be derived, not all of which are testable. Instead, it is enough to focus on testing the surprising consequences, which, when enough are verified, make it possible to reliably predict perceptual performance in unstudied domains. Past and recent work has built on the Bayesian perspective to advance a number of what we might call “deep principles” applicable to human perception.

- 1) The visual system seeks codes which minimize redundancy in the input (Barlow, 1959; Olshausen and Field, 1996; Atick and Redlich, 1992; Bell and Sejnowski, 1997). This principle exists in various forms, such as MDL encoding, minimax entropy (Zhu et al., 1997), principal components analysis (Bossomaier and Snyder, 1986) and independent components analysis (Bell and Sejnowski, 1997).
- 2) Given equally likely causes of an image, the visual system chooses the model with the least number of assumptions. In this sense, quantitative versions of the Gestalt principle of simplicity (e.g. via MDL realization of Occam’s razor) apply as a principle to resolve ambiguity (Restle, 1982; Leeuwenberg, 1969). The pattern inference theory distinctive is that it has the (yet to be obtained) goal of deriving the rules of simplicity from density models based on ensembles of natural image (e.g. (Zhu, 1999)).
- 3) The visual system actively acquires new information by maximizing the expected utility or minimizing entropy of the information for the task (Amit and Geman, 1997). This principle has been applied to an ideal observer model of human reading (Legge et al., 1997).
- 4) Perceptual decisions are confidence-driven. This requires that computations take into account both estimates and the degree of uncertainty in those estimates. Evidence that human perception does this comes from studies on cue integration (Landy et al., 1995), orientation from texture discussed above (Knill, 1998a), motion perception, discussed below (Weiss and Adelson, 1998), and visual motor control (Wolpert et al., 1995).
- 5) Perception’s model of the image measurement should be robust with respect to variations in the secondary variables. We noted above that the general viewpoint principle is a consequence of viewpoint being a secondary variable (Freeman, 1994), and that ambiguity in depth from shadows can be resolved by treating illumination direction as secondary.
- 6) The visual system predictively models its behavioral outcomes. Until recently, the Bayesian approach to perception has been largely static; however, Bayesian techniques can be used to model both learning (Jordan, 1998) and time-variant processes (Dean, 1988; Barker et al., 1995). (*Myth 4: Bayes lacks dynamics.*) For example, the Kalman filter provides a good account of kinematics of visual control of human reach (Wolpert et al., 1995). Consistent with the probability computation theme of this chapter, the Kalman filter goes beyond estimates of central tendency, and estimates both the mean and variance of control parameters.
- 7) The visual system performs ideal inference given its limitations in representing image data, but only for a limited number of tasks (Schrater and Kersten, 2000). In the next section, we discuss using this principle to develop models of ideal performance as a default hypotheses. It is essentially a statement that the visual system should be optimally adapted to perform certain visual tasks relevant to the observer’s needs.

For Bayesian theory construction to be useful, we must show that the theories admit experimental testing. In the next section we discuss practical aspects of testing pattern theoretic hypotheses at several levels of specificity. In particular, we return to principle (7).

6.2 Ideal observers and human scene inference

How do we formulate and test theories of human behavioral function within a pattern inference theory framework? In psychophysical experiments, one can: a) test at the constraint level—what information does human vision avail itself of?, or; b) test at the mechanism level—what neural subsystem can account for performance? Pattern inference theory is of primary relevance to hypotheses testable at the former level. Tests of human perception can be based on hypotheses regarding constraints contained in: the two components of the generative model, 1) the prior $p(S)$ and 2) the likelihood $p(I|S)$; 3) the image model $p(I)$; or 4) the posterior $p(S|I)$. A distinction based on the source of a constraint serves to clarify the otherwise confusing idea of “cue” which muddles scene and image constraints (Knill et al., 1996b). For example, the “occlusion cue” is sometimes defined in terms of “overlapping surfaces”, and sometimes as a “T-junction” in the image contours. But surface occlusion is the *causal source* of a “T-junction”. (*Myth 5: Identifying “Bayesian constraints” provides no advantage over identifying traditional “cues”.*)

1) *The prior.* The well-known “light from above” assumption in shape-from-shading is an example of an hypothesis expressed solely in terms of a prior distribution on a scene variable, light source direction. Given that primary lighting for most of human history has been the sun, a prior bias on lighting from above is an example of a prediction which could be generated by a study of the natural distribution of scene variables, which can be quantitatively documented using density estimation. A fruitful first pass could be a more widespread use of principal components analysis as a way of seeking economical density models. Indeed, empirical measurements of the distribution of spectral reflectance functions of natural surfaces have shown that the set of naturally occurring spectral reflectance functions can be well-modeled as linear combinations of three basis functions (Maloney and Wandell, 1986). When restricted to natural illumination conditions, this result supplies an especially simple interpretation of trichromacy: three spectral measurements are usually enough to determine spectral reflectance. Earlier we noted research on prior models for facial surfaces (Atick et al., 1996; Vetter and Troje, 1997). In a different example, an observer’s assessment of the 3-D position of a moving ball is affected by moving cast shadow information. The observer’s data can be qualitatively described in terms of a prior “stationary light source” constraint (Knill et al., 1996a). The subjective biases in the perception of shape from line contours have been studied by Mamassian and Landy (1998) (Mamassian and Landy, 1998). An interesting problem for the future will be to relate these subjective priors to ones discovered objectively through density estimation (Zhu, 1999).

2) *The likelihood term.* The independent variables in a psychophysical experiment can be specified in terms of the scene or image variables, or in the language of perceptual psychology, in terms of the distal or the proximal stimulus. Even if one doesn’t have an ideal observer model, it is still possible to manipulate the scene variables in the generative model to test hypotheses at these levels, if one has some way to account for changes in performance due to changes in image information. For example, object recognition must deal with variations in both viewpoint and illumination. View-based theories of object recognition rest on experiments showing that human vision doesn’t compensate for all view variations with equal facility (Tarr and Bülthoff, 1995). Scale changes are handled with less sensitivity to view familiarity than either rotations in the image or rotations in depth. However, the degree to which the human visual system is view-dependent will require developing ideal observer models for object

recognition, because part of performance variation due to viewpoint can be due to the informativeness of the viewpoint for the recognition task. In fact, Liu et al. (1996) showed that human observers are more efficient than simple view-based algorithms at recognizing simple wire objects (Liu et al., 1995).

Object recognition must also compensate for illumination variations. The fact, mentioned above, that under fairly general conditions, the space of images generated by an object under fixed view is a cone in image space makes predictions regarding how object recognition should generalize under illumination change, as well as the discriminability of two objects with distinct illumination cones (Belhumeur and Kriegman, 1996; Tarr et al., 1998).

3) *The image density.* Given a model, $p(I; \Lambda)$, of an image ensemble in a domain Λ (i.e. natural image prior, or more specific texture priors, such as “fur”), one can test how well the human visual system is “tuned” to the statistical structure specified by the parameters Λ . An example of this approach is the ideal observer analysis of human discrimination of Gaussian textures with $1/f^\alpha$ spatial frequency spectra (Knill et al., 1990). Current theoretical work on non-Gaussian texture modeling (e.g. Minimax entropy discussed above) is providing richer models of natural images that provide testable psychophysical hypotheses (Buccigrossi and Simoncelli, 1997; Zhu et al., 1997; Ruderman and Bialek, 1994). More domain-specific density models, e.g. using PCA, have been used to model face variation in the image domain (Sirovich and Kirby, 1987; Turk and Pentland, 1991), and have motivated psychological theories (Valentin et al., 1997).

4) *The posterior.* A full quantitative model (the grand challenge) requires tests at the level of the posterior. A statistical theory of visual inference plays two roles, it *normalizes* performance and *models* perception. An ideal observer model, which bases its performance on $P(S|I)$, provides the benchmark to normalize human performance relative to the information available for the task (Barlow, 1962). The importance of this normative measure cannot be overstated. Without carefully assessing how the information changes across experimental conditions, the mechanisms underlying changes in performance become nearly impossible to determine. In fact, normalizing human performance with respect to the available information can lead to the opposite conclusions from those based on the unnormalized performance (Eagle and Blake, 1995).

But in what sense does an ideal observer serve a modeling function? The fact that perception enables successful behavior has a non-trivial impact on the principles required to understand perception. In this regard, pattern inference theory is sympathetic with one aspect of the ecological approach to perception, namely that theories of visual perception cannot be built in isolation from functional visual tasks. If this is indeed the case, our grand challenge is unavoidably grand. This raises a dilemma for scientific methodology. If we have to worry about the large set of variables involved in normal perception, how can we manage controlled experimental tests of the theory? We believe the answer is to use the ideal observer as the default experimental hypothesis—in other words, first test whether human vision utilizes the information available for the task optimally. Of course, in general it won’t. However, because the ideal starts off (at least in principle) with no free parameters, a sub-optimal theory can still achieve economy through modification of the ideal theory through the frugal introduction of free parameters. Further, any parameters introduced should be related to some biologically (or ecologically) relevant limitation on processing. This idea is very much in the spirit of sequential ideal observer analysis of photon and biological limits to resolution (Geisler, 1989). In the domain of surface perception, Knill has shown that human discrimination of surface slant improves with slant—a behavior that can be predicted from an ideal observer analysis of the information (Knill, 1998a).

As mentioned earlier, the true test of a quantitative framework such as pattern inference theory is its

ability to generate economical and predictive theories. A rather striking recent success story is Weiss and Adelson's Bayesian theory of motion perception in which they tackled the problem of combining local motion measurements. Previously, distinct models (often with distinct hypothesized physiological realizations) have been proposed for various classes of motion effects. Weiss and Adelson were able to account for a wide and diverse range of human perceptual results with only one free parameter, the uncertainty in the local motion measurement (Weiss and Adelson, 1998). By assuming that the visual system takes into account the uncertainties in local motion measurements, rather than just the motion estimates (i.e. confidence-driven processing), and assuming simple priors for slower speeds and smooth velocity fields, the MAP estimate from the resulting posterior modeled: the barber-pole effect, the biases in the direction and speed of moving plaids due to differences in component orientation and contrast, the wagon-wheel effect, non-rigid contour perception, as well as others.

A crucial aspect to the success of such a program is to choose an information processing task for which the human visual system is well-suited. Our assumption is that if this is done, the ideal performance will be a good first-approximation to a model for human performance. This argument is similar to those made with regard to adaptability in cognitive tasks (Cosmides, 1989; Murray, 1987; Gigerenzer, 1998). Pattern inference theory provides the language to express experimentally testable theories in a way analogous to calculus being useful for expressing quantitative theories in science generally. As such it is a mathematical theory, not a falsifiable experimental theory. However, pattern inference theory provides the means to generate testable scientific theories of perception with few free parameters; such theories should be more easily falsifiable than, for example connectionist theories with lots of free parameters (*Myth 6: Bayesian theories are not falsifiable*).

7 Does the brain compute probabilities?

7.1 Computing probabilities is not probabilistic computing.

Shortly before his death in 1957, John von Neumann wrote "The language of the brain is not the language of mathematics" (Von Neumann, 1958). He went on to predict that brain theory would eventually come to resemble the physics of statistical mechanics and thermodynamics. His argument was based on the apparent imprecision in neural pulse-train coding, rather than an underlying computational function for stochastic processing elements. Nevertheless, Von Neumann would no doubt have been intrigued by the algorithms, developed since, that do in fact rely on stochastic processing. But it is one thing to say that brain processing is limited by neural noise, or that it uses noise to compute, and quite another to state that the brain computes probabilities. An information processing system can compute probabilities quite deterministically (see example below).

The main purpose of this chapter is to argue that the perceptual system must necessarily do statistical inference, and thus compute probabilities. The degree of sophistication of such processes and the brain mechanisms are open questions. For example, a standard assumption in vision has been that probability computations are limited to computing central tendencies of distributions, which produce estimates of quantities of interest for decisions. But this leads to a premature commitment by discarding the information about the uncertainty in the estimate. In addition to the empirical work discussed earlier, it has also been shown from work in Bayes nets that significant improvements in computational convergence time are obtained when probabilities are propagated (e.g. (Weiss, 1997)), thus preserving information about the degree of uncertainty.

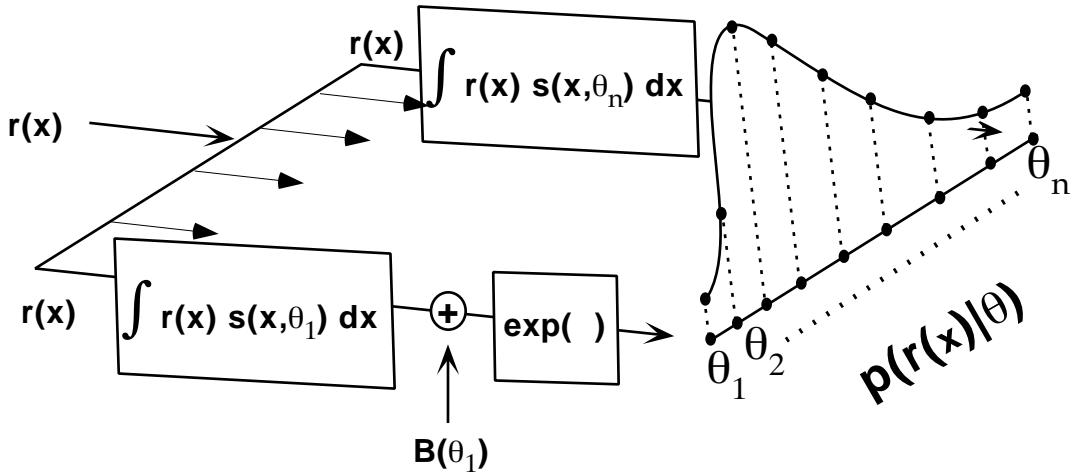


Figure 6: Computing probabilities in visual cortex. Probabilities can be computed across ‘labeled line’ maps of neurons. The diagram shows an example of how a set of linear receptive fields can compute a posterior distribution. Assume the visual signal $r(\vec{x})$ can be decomposed into a sum of basis signals which vary according to some parameter θ plus some gaussian image noise: $r(\vec{x}) = \sum_i s(\vec{x}, \theta_i) + N$. For instance, the image could be decomposed into a sum oriented Gabor patches. The sampled (unnormalized) posterior probability of the parameter θ can be computed using an array of linear filters, each corresponding to a particular value of θ . The inner product of the signal with the receptive field produces the log likelihood for the presence of the component $s(\vec{x}, \theta_i)$. The log prior probability and needed bias are lumped as the additive constant $B(\theta)$, and the result is mapped from log likelihoods to probabilities by passing the results through a point-wise accelerating $\exp()$ non-linearity.

7.2 How could probability densities be represented in the brain?

Although Bayesian theories require the computation of probabilities, they do not necessarily require probabilities to be explicitly computed. For example, the knowledge regarding probability densities could be *implicitly* encoded in neural non-linearities and neural weights. For example, the form of the photoreceptor transducer can be interpreted as the result of mapping a non-uniform distribution of contrasts to a uniformly distributed response variable—i.e. histogram equalization (Laughlin, 1981). Knowledge of the image probability density function is implicit in the photoreceptor in the sense that the density is the derivative of its transducer function.¹⁴

At a higher level, for a classification task the visual system need only deterministically map the image data onto an appropriate decision variable space and then choose the category boundaries, neither of which require probabilities to be explicitly computed.

On the other hand, the visual system may *explicitly* compute probabilities in terms of population codes across a variable of interest. The basic idea is straightforward: the firing rate of a labeled line code for an image or scene variable is proportional to the posterior probability for the variable. Population codes occur in the coding of motor planning for eye movements in superior colliculus (Lee et al., 1988) and for reaching (Georgopoulos et al., 1989) in motor cortex, and they have been proposed for the representation

¹⁴This is a consequence of the density mapping theorem: $p_y(y) = \int \delta(y - f(x))f^{-1}(x)p_x(x)dx$ over each monotonic part of f .

of sensory information such as a population code for image velocities (Simoncelli, 1993; Schrater, 1998). Population codes hold the promise of being able to represent probabilities for computation. For instance, the uncertainty can be represented as the entropy over the ensemble or spread of activity across the population, and multiple values can be represented as multi-modal distributions of activity. Probability information can be transmitted from one functional group of neurons to the next by transforming the distributions between parameter spaces using the density mapping theorem. In contrast, a system that performs “estimation propagation” would summarize the population first (e.g. mean or mode) into an estimate of state, and then only propagate the estimate. Of course, making discrete decisions is one of the main types of task discussed above, and at some point the probability distribution could be collapsed to one decision variable.

The connection between population codes and probability distributions goes back to early work in communication engineering, where it was shown that an array of linear filters could produce a *sampled* log likelihood function (Van Trees, 1971). As an example, assume the input visual signal $r(\vec{x})$, is a sum of components of interest, $s(\vec{x}, \theta_i)$, plus Gaussian noise N : $r(\vec{x}) = \sum_i s(\vec{x}, \theta_i) + N$. The components could be oriented edge segments, in which case $s(\vec{x}, \theta_i)$ are prototypical edge images, parametrized by the orientation θ_i . Given the input signal $r(\vec{x})$, it is straightforward to show that the likelihood function for the signal given an orientation is given by $\log p(r(\vec{x})|\theta_i) = \int_{\vec{x}} r(\vec{x})s(\vec{x}, \theta_i)d\vec{x}$. This likelihood can be converted into an unnormalized posterior probability by adding a constant equal to the log prior probability θ_i , and running the outputs through an accelerating exponential non-linearity (see figure 6). The limitation to sampled likelihood functions can be overcome by using “steerable filters” (Freeman, 1992), which allow the likelihood values between the samples to be computed as weighted sums of the samples. Given the linear receptive fields in visual area V1, this simple example may have an analog in the brain. The key idea is that a simple filtering operation yields a neuron whose firing rate can represent a likelihood. For local image velocity estimation, Simoncelli (Simoncelli et al., 1991; Simoncelli, 1993) showed how the posterior distribution for local image velocity could be computed in terms of simple combinations of the outputs of spatial and temporal derivatives of Gaussian filters. More recently, Schrater (Schrater, 1998) has shown how the likelihood for image velocity could be computed by the weighted sum of a set of units similar to V1 complex cells.

Several authors have proposed more general ways in which probability distributions could be represented by a population of neurons (Anderson, 1994; Sanger, 1996; Zemel, 1997, 1998). The basic idea is that a population of neurons which are tuned to some parameter can act like a kernel density estimate, in which a probability density is approximated as a combination of simpler densities. To illustrate, assume we have a posterior distribution $p(x|D)$ of some scene variable x (like position), given a set of image measurements D , and a set of receptive fields $f_i(x)$. Then the firing rates of the neurons will be given by the projection $r_i = \int_x p(x|D)f_i(x)dx$. Unlike the previous examples, the firing rate computed this way is not explicitly proportional to a probability or likelihood. Instead, the posterior is coded implicitly by the firing rates. To explicitly compute the posterior from the r_i , a fixed set of interpolation kernels $\phi_i(x)$ is used to invert the projection. Zemel et al. (Zemel, 1998) discuss two similar schemes to do the encoding and decoding of posterior distributions from firing rates. The number of ways of encoding posterior distributions by similar methods is limitless, and whether or not the brain uses a particular scheme of this sort is an intriguing problem for the future.

8 Summary

We have argued that probability computation by the visual system is a necessary consequence of nature’s visual signal patterns being inherently statistical. The origins of this perspective on perception began with the development of signal detection theory, and in particular, with ideal observer analysis. The basic operations of probability theory provide the means to model information for a task, and decision theory provides tools to model task requirements. The application of these tools to natural pattern understanding falls in the domain of what we have referred to as pattern inference theory—the combination of Bayesian decision theory and pattern theory. Pattern inference theory is clearly more powerful than any specific experimentally testable theory of human perception. However, it provides a sufficiently rich language to develop theories of natural perceptual function. In addition to reviewing a number of principles that fall out of the Bayesian formalism, we highlighted two relatively new principles: 1) A Bayesian principle of least commitment, in which one propagates probabilities, rather than estimates, thereby weighting evidence according to reliability; 2) A Bayesian principle of modularity, in which Bayes nets show how statistical structure and task determine modularity.

9 Acknowledgments

We thank Larry Maloney for providing exceptionally thoughtful and constructive criticisms on the first draft of this paper. This research was supported by NSF SBR-9631682 and NIH RO1 EY11507-001.

References

1. Amit, Y., & Geman, D. (1997). Shape quantization and recognition with random trees. *Neural Computation*, **9**(7), 1545–1588.
2. Anderson, C. H. (1994). Basic elements of biological computational systems. *International Journal of Modern Physics C*, **5**(2), 135–137.
3. Atick, J. J., Griffin, P. A., & Redlich, A. N. (1996). Statistical Approach to Shape from Shading: Reconstruction of Three-Dimensional Face Surfaces from Single Two-Dimensional Images. *Neural Computation*, **8**(6), 1321–1340.
4. Atick, J. J., & Redlich, A. N. (1992). What does the retina know about natural scenes? *Neural Computation*, **4**(2), 196–210.
5. Barker, A., Brown, D., & Martin, W. (1995). Bayesian estimation and the Kalman filter. *Computers Math. Applic.*, **30**(10), 55–77.
6. Barlow, H. (1959). Sensory mechanisms, the reduction of redundancy, and intelligence. In *Proceedings of the symposium on the mechanization of thought processes*, National Physical Laboratory. HMSO, London.
7. Barlow, H. (1962). A method of determining the overall quantum efficiency of visual discriminations. *J. Physiol. (Lond.)*, **160**, 155–168.
8. Barlow, H. (1981). Critical Limiting Factors in the Design of the Eye and Visual Cortex. *Proc. Roy. Soc. Lond. B*, **212**, 1–34.
9. Belhumeur, P., & Kriegman, D. (1996). What is the set of images of an object under all possible lighting conditions? In *IEEE Conf. on Computer Vision and Pattern Recognition*, pages 270–277, San Francisco, CA.
10. Bell, A. J., & Sejnowski, T. J. (1997). The "independent components" of natural scenes are edge filters. *Vision Res*, **37**(23), 3327–38.
11. Berger, J. (1985). *Statistical Decision Theory and Bayesian Analysis*: Springer.
12. Bloj, M. G., Kersten, D., & Hurlbert, A. C. (1999). Perception of three-dimensional shape influences colour perception through mutual illumination. *Nature*, **402**(6764), 877–879.
13. Bobick, A. (1997). Movement, activity and action: the role of knowledge in the perception of motion. *Philosophical Transactions of the Royal Society B*, **352**(1358), 1257–1265.
14. Bobick, A., & Richards, W. (1986). Classifying Objects from Visual Information. Technical Report A.I. Memo No. 879, Artificial Intelligence Laboratory Massachusetts Institute of Technology.
15. Bookstein, F. L. (1978). *The Measurement of Biological Shape & Shape Change*: Springer-Verlag, New York Incorporated.
16. Bookstein, F. L. (1997). *Morphometric Tools for Landmark Data: Geometry and Biology*: Cambridge University Pres.

17. Bossomaier, T., & Snyder, A. (1986). Why Spatial Frequency Processing in the Visual Cortex? *Vision Research*, **26**(8), 1307–1309.
18. Brainard, D. H., & Freeman, W. T. (1997). Bayesian color constancy. *J Opt Soc Am A*, **14**(7), 1393–411.
19. Buccigrossi, R., & Simoncelli, E. (1997). Progressive Wavelet Image Coding Based on a Conditional Probability Model. In *ICASSP*, Munich, Germany.
20. Cosmides, L. (1989). The logic of social exchange: Has natural selection shaped how humans reason? Studies with the Wason selection task. *Cognition*, **31**, 187–276.
21. Cover, T. M., & Joy, A. T. (1991). *Elements of Information Theory*. Wiley Series in Telecommunications. New York: John Wiley & Sons, Inc.
22. Craik, K. J. W. (1943). *The Nature of Explanation*. Cambridge: Cambridge University Press.
23. Dayan, P., Hinton, G. E., Neal, R. M., & Zemel, R. S. (1995). The Helmholtz Machine. *Neural Computation*, **7**(5), 889–904.
24. Dean, T. & Kanazawa, K. (1988). Probabilistic temporal reasoning. In *AAAI-88*, pages 524–528.
25. Duda, R., & Hart, P. (1973). *Pattern classification and scene analysis*. New York.: John Wiley & Sons.
26. Eagle, R. A., & Blake, A. (1995). Two-dimensional constraints on three-dimensional structure from motion tasks. *Vision Res*, **35**(20), 2927–41.
27. Egan, J. P. (1975). *Signal detection theory and ROC-analysis*. Academic Press series in cognition and perception. New York: Academic Press.
28. Epstein, R., Hallinan, P., & Yuille, A. (1995). $5 \pm$ Eigenimages Suffice: An Empirical Investigation of Low-Dimensional Lighting Models. In *IEEE Workshop on Physics-Based Modeling in Computer Vision*, pages 108–116, Boston, MA.
29. Field, D. J. (1987). Relations between the statistics of natural images and the response properties of cortical cells. *Journal of the Optical Society of America A*, **4**(12), 2379–2394.
30. Freeman, W. T. (1992). Steerable Filter and Local Analysis of Image Structure. Technical Report 190, Massachusetts Institute of Technology.
31. Freeman, W. T. (1994). The generic viewpoint assumption in a framework for visual perception. *Nature*, **368**(7 April 1994), 542–545.
32. Freeman, W. T., & Pasztor, E. C. (1999). Learning to estimate scenes from images. In M. S. Kearns, S. A. S., & Cohn, D. A., editors, *Adv. Neural Information Processing Systems 11*. MIT Press, Cambridge MA.
33. Frey, B. J. (1998). *Graphical Models for Machine Learning and Digital Communication*. Adaptive Computation and Machine Learning series. A Bradford Book. Cambridge, Massachusetts: MIT Press.
34. Geisler, W. (1989). Sequential Ideal-Observer analysis of visual discriminations. *Psychological Review*, **96**(2), 267–314.

35. Georgopoulos, A., Lurito, J., Petrides, M., Schwartz, A., & Massey, J. (1989). Mental Rotation of the Neuronal Population Vector. *Science*, **243**, 234–236.
36. Gibson, J. J. (1966). *The Senses Considered as Perceptual Systems*. Boston, MA: Houghton Mifflin.
37. Gigerenzer, G. (1998). Ecological intelligence: An adaptation for frequencies. In Cummins, D. D., & Allen, C., editors, *The Evolution of Mind*. Oxford University Press, Oxford.
38. Green, D. M., & Swets, J. A. (1974). *Signal Detection Theory and Psychophysics*. Huntington, New York: Robert E. Krieger Publishing Company.
39. Gregory, R. (1980). Perceptions as hypotheses. *Philosophical Transactions of the Royal Society*, **B 290**, 181–197.
40. Grenander, U. (1950). Stochastic processes and statistical inference. *Arkiv. Mathematik*, **1**(17), 195.
41. Grenander, U. (1993). *General Pattern Theory*. Oxford: Oxford University Press.
42. Grenander, U. (1996). *Elements of Pattern theory*. Baltimore, MD: Johns Hopkins University Press.
43. Grenander, U., Chow, Y., & Keenan, D. M. (1991). *Hands. A Pattern Theoretic Study of Biological Shapes*. New York: Springer.
44. Helmholtz, H. v., & Southall, J. P. C. (1924). *Helmholtz's treatise on physiological optics*. Rochester, N.Y.: The Optical Society of America.
45. Hinton, G., & Ghahramani, Z. (1997). Generative models for discovering sparse distributed representations. *The Philosophical Transactions of the Royal Society*, **352**(1358), 1177–1190.
46. Jaynes, E. T. (1957). Information theory and statistical mechanics. *Physical Review*, **106**, 620–630.
47. Jordan, M. (1998). *Learning in Graphical Models*. Cambridge, MA: MIT Press.
48. Kendall, D. (1989). A survey of the statistical theory of shape. *Statistical Science*, **4**(2), 87–120.
49. Kersten, D. (1984). Spatial summation in visual noise. *Vision Research*, **24**, 1977–1990.
50. Kersten, D. (1990). Statistical limits to image understanding. In Blakemore, C., editor, *Vision: Coding and Efficiency*, pages 32–44. Cambridge University Press, Cambridge, UK.
51. Kersten, D. (1997). Perceptual categories for spatial layout. *Philosophical Transactions of the Royal Society B*, **352**(1358), 1155–1163.
52. Kersten, D. (1999). High-level vision as statistical inference. In Gazzaniga, M. S., editor, *The New Cognitive Neurosciences – 2nd Edition*, pages 353–363. MIT Press, Cambridge, MA.
53. Kersten, D., O'Toole, A., Sereno, M., Knill, D. C., & Anderson, J. (1987). Associative learning of scene parameters from images. *Appl. Opt.*, **26**, 4999–5006.
54. Kersten, D. J. (1991). Transparency and the Cooperative Computation of Scene Attributes. In Landy, M., & Movshon, A., editors, *Computational Models of Visual Processing*, pages 209–228. M.I.T. Press, Cambridge, Massachusetts.
55. Knill, D., Field, D., & Kersten, D. (1990). Human discrimination of fractal images. *7*, 1113–1123.

56. Knill, D., & Richards, W. (1996). *Perception as Bayesian Inference*. Cambridge: Cambridge University Press.
57. Knill, D. C. (1998a). Discrimination of planar surface slant from texture: human and ideal observers compared. *Vision Res*, **38**(11), 1683–711.
58. Knill, D. C. (1998b). Surface orientation from texture: ideal observers, generic observers and the information content of texture cues. *Vision Research*, **38**(11), 1655–1682.
59. Knill, D. C., Kersten, D., & Mamassian, P. (1996a). The Bayesian Framework for visual information processing: implications for psychophysics. In D.C., K., & W., R., editors, *Perception as Bayesian Inference*, pages 239–286, Chap. 5. Cambridge University Press.
60. Knill, D. C., Kersten, D., & Yuille, A. (1996b). A Bayesian Formulation of Visual Perception. In D.C., K., & W., R., editors, *Perception as Bayesian Inference*, page Chap. 0. Cambridge University Press.
61. Landy, M. S., Maloney, L. T., Johnston, E. B., & Young, M. J. (1995). Measurement and modeling of depth cue combination: In defense of weak fusion. *Vision Research*, **35**, 389–412.
62. Laughlin, S. (1981). A simple coding procedure enhances a neuron's information capacity. *Z. Naturforsch.*
63. Lee, C., Rohrer, W. H., & Sparks, D. L. (1988). Population coding of saccadic eye movements by neurons in the superior colliculus. *Nature*, **332**(6162), 357–60.
64. Leeuwenberg, E. (1969). Quantitative specification of information in sequential patterns. *Psychological Review*, **76**, 216–220.
65. Legge, G. E., Klitz, T. S., & Tjan, B. S. (1997). Mr. Chips: an ideal-observer model of reading. *Psych. Review*, **104**(3), 524–53.
66. Liu, Z., Knill, D. C., & Kersten, D. (1995). Object Classification for Human and Ideal Observers. *Vision Research*, **35**(4), 549–568.
67. MacKay, D. J. C. (1992). Bayesian interpolation. *Neural Computation*, **4**(3), 415–447.
68. Maloney, L., & Wandell, B. (1986). Color Constancy: A Method for Recovering Surface Spectral Reflectance. *Journal of the Optical Society America*, **3**, 29–33.
69. Mamassian, P., & Landy, M. S. (1998). Observer biases in the 3D interpretation of line drawings. *Vision Res*, **38**(18), 2817–32.
70. Marr, D. (1982). *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information*. San Francisco, CA: W.H. Freeman and Company.
71. Mumford, D. (1992). On the computational architecture of the neocortex. II. The role of cortico-cortical loops. *Biological Cybernetics*, **66**, 241–251.
72. Mumford, D. (1994). Neuronal architectures for pattern-theoretic problems. In Koch, C., & Davis, J. L., editors, *Large-Scale Neuronal Theories of the Brain*, pages 125–152. MIT Press, Cambridge, MA.

73. Mumford, D. (1996). Pattern theory: A unifying perspective. In Knill, D., & W., R., editors, *Perception as Bayesian Inference*, page Chapter 2. Cambridge University Press, Cambridge.
74. Murray, G. . (1987). *Cognition as Intuitive Statistics*: Erlbaum.
75. Nakayama, K., & Shimojo, S. (1992). Experiencing and perceiving visual surfaces. *Science*, **257**, 1357–1363.
76. Neyman, J., & Pearson, E. (1933). On the problem of the most efficient tests of statistical hypotheses. *Phil. Trans. Roy. Soc. London, Series A*, page 289.
77. Olshausen, B., & Field, D. (1996). Emergence of simple-cell receptive field properties by learning a sparse code for natural images. *Nature*, **381**, 607–609.
78. Pearl, J. (1988). *Probabilistic Reasoning in Intelligent Systems*. San Mateo, CA: Morgan Kaufmann Publishers Inc.
79. Pelli, D. G. (1990). The quantum efficiency of vision. In Blakemore, C., editor, *Vision: Coding and Efficiency*. Cambridge University Press, Cambridge.
80. Peterson, W., Birdsall, T., & Fox, W. (1954). The theory of signal detectability. *Trans. IRE Professional Group on Information Theory*, **PGIT-4**, 171–212.
81. Poggio, T., & Girosi, F. (1990). Regularization algorithms for learning that are equivalent to multilayer networks. *Science*, **247**, 978–982.
82. Poggio, T., Torre, V., & Koch, C. (1985). Computational vision and regularization theory. *Nature*, **317**, 314–319.
83. Rao, R. P., & Ballard, D. H. (1999). Predictive coding in the visual cortex: a functional interpretation of some extra-classical receptive-field effects [see comments]. *Nat Neurosci*, **2**(1), 79–87.
84. Restle, F. (1982). Coding theory as an integration of Gestalt psychology and information processing theory. In Beck, J., editor, *Organization and Representation in Perception*, pages 31–56. Erlbaum, Hillsdale, NJ.
85. Rice, S. O. (1944). Mathematical Analysis of Random Noise. *Bell System Technical Journal*, **23**, 282–332.
86. Ripley, B. (1996). *Pattern Recognition and Neural Networks*. Cambridge, UK: Cambridge University Press.
87. Rock, I. (1983). *The Logic of Perception*. A Bradford Book. Cambridge, Massachusetts: M.I.T. Press.
88. Rosch, E., Mervis, C., Gray, W., Johnson, D., & Boyes-Braem, P. (1976). Basic objects in natural categories. *Cognitive Psychology*, **8**, 382–439.
89. Ruderman, D., & Bialek, W. (1994). Statistics of Natural Images: Scaling in the Woods. *Physical Review Letters*, **73**, Number **6**(8 August 1994), 814–817.
90. Sanger, T. D. (1996). Probability density estimation for the interpretation of neural population codes. *Journal of Neurophysiology*, **76**(4), 2790–2793.

91. Schrater, P. (1998). *Local motion detection: Comparison of human and model observers*. Ph.d., University of Pennsylvania.
92. Schrater, P. R., & Kersten, D. (2000). The role of task specification in optimal cue integration. *International Journal of Computer Vision*, **40**(1), 71–89.
93. Schwarz, G. (1978). Estimating the dimension of a model. *Ann. Stat.*, **6**, 461–463.
94. Shannon, C. E., & Weaver, W. (1949). *The mathematical theory of communication*. Champaign, IL: U. Illinois Press.
95. Simoncelli, E. P. (1993). *Distributed Analysis and Representation of Visual Motion*. Ph.d., Massachusetts Institute of Technology, Department of Electrical Engineering and Computer Science.
96. Simoncelli, E. P. (1997). Statistical Models for Images: Compression, Restoration and Synthesis. In *Proc. 31st Asilomar Conference on Signals, Systems and Computers*, Pacific Grove, CA. (c) IEEE Signal Processing Society.
97. Simoncelli, E. P., Adelson, E. H., & Heeger, D. J. (1991). Probability Distributions of Optical Flow. In *IEEE Conf on Computer Vision and Pattern Recognition*, Maui, Hawaii.
98. Simoncelli, E. P. & Portilla, J. (1998). Texture Characterization via Joint Statistics of Wavelet Coefficient Magnitudes. In *5th IEEE International Conference on Image Processing*, Chicago, IL.
99. Sirovich, L., & Kirby, M. (1987). Low-dimensional procedure for the characterization of human faces. *JOSA*, **4**(3), 519–524.
100. Swets, J. A. (1988). Measuring the accuracy of diagnostic systems. *Science*, **240**(4857), 1285–93.
101. Tarr, M., & Bülthoff, H. (1995). Is human object recognition better described by geon-structural-descriptions or by multiple-views? *Journal of Experimental Psychology: Human Perception and Performance*, **21**(6), 1494–1505.
102. Tarr, M. J., Kersten, D., & Bulthoff, H. H. (1998). Why the visual recognition system might encode the effects of illumination. *Vision Res*, **38**(15-16), 2259–75.
103. Turk, M., & Pentland, A. (1991). Eigenfaces for recognition. *Journal of Cognitive Neuroscience*, **3**(1), ??
104. Valentin, D., Abdi, H., Edelman, B., & O'Toole, A. J. (1997). Principal Component and Neural Network Analyses of Face Images: What Can Be Generalized in Gender Classification? *J Math Psychol*, **41**(4), 398–413.
105. Van Trees, H. L. (1968). *Detection, Estimation and Modulation Theory. Part I.*, volume 1. New York: John Wiley and Sons.
106. Van Trees, H. L. (1971). *Detection, Estimation and Modulation Theory. Part III.*, volume 3. New York: John Wiley and Sons.
107. Vapnik, V. (1995). *The nature of statistical learning*. New York: Springer-Verlag.
108. Vetter, T., & Troje, N. (1997). Separation of texture and shape in images of faces for image coding and synthesis. *Journal of the Optical Society of America A*, **14**(9), 2152–2161.

109. Von Neumann, J. (1958). *The computer and the brain*. New Haven,: Yale University Press.
110. Wald, A. (1939). Contributions to the of statistical estimation and testing hypotheses. *Ann. Math. Stat.*, **10**, 299–326.
111. Wald, A. (1950). *Statistical Decision Functions*. New York: John Wiley & Sons.
112. Weiss, Y. (1997). Interpreting images by propagating Bayesian beliefs. In M.C. Mozer, M. J., & Petsche, T., editors, *Advances in Neural Information Processing Systems 9*, pages 908–915. MIT Press, Cambridge MA.
113. Weiss, Y., & Adelson, E. H. (1998). Slow and smooth: a Bayesian theory for the combination of local motion signals in human vision. Technical Report A.I. Memo No. 1624, M.I.T.
114. Wolpert, D. M., Ghahramani, Z., & Jordan, M. I. (1995). An Internal Model for Sensorimotor Integration. *Science*, **269**(29 September), 1880–1882.
115. Yuille, A. (1991). Deformable templates for face recognitoin. *Journal of Cognitive Neuroscience*, **3**(1), 59–70.
116. Yuille, A. L., & Blthoff, H. H. (1996). Bayesian decision theory and psychophysics. In D.C., K., & W., R., editors, *Perception as Bayesian Inference*. Cambridge University Press, Cambridge, U.K.
117. Yuille, A. L., Coughlan, J. M., & Kersten, D. (1998). Computational Vision: Principles of Perceptual Inference.
118. Zemel, R. S. (1997). Combining probabilistic population codes. In *International Joint Conference on Artificial Intelligence*, Denver, CO. Morgan Kaufmann.
119. Zemel, R. S. (1998). Probabilistic interpretation of population codes. *Neural Computation*, **10**(2), 403–430.
120. Zhu, S. (1999). Embedding Gestalt Laws in Markov Random Fields. *IEEE Trans. Pattern Analysis and Machine Intelligence.*, **21**(11).
121. Zhu, S., & Mumford, D. (1998). GRADE: A framework for pattern synthesis, denoising, image enhancement, and clutter removal. In *Proceedings of International Conference on Computer Vision*, Bombay, India. Morgan Kaufmann.
122. Zhu, S. C., Wu, Y., & Mumford, D. (1997). Minimax Entropy Principle and Its Applications to Texture Modeling. *Neural Computation*, **9**(8), 1627–1660.